

Neural Networks


Vlad Gladkikh

IBS CMCM




Borophene with Large Holes

Yong Wang, Yunjae Park, Lu Qiu, Izaac Mitchell, and Feng Ding*

J. Phys. Chem. Lett. 2020, 11, 6235–6241



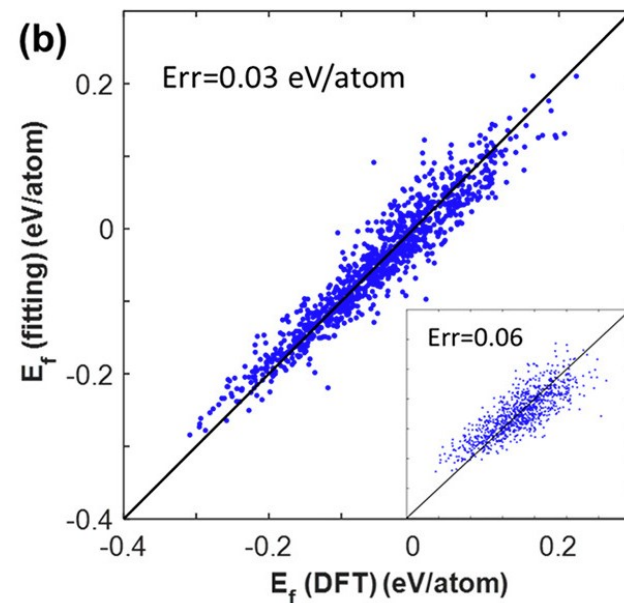
Neurons and the brain



humans don't
need features

Lecture 20 of the Introductory Applied Machine Learning (IAML)
course at the University of Edinburgh, taught by Victor Lavrenko

https://www.youtube.com/playlist?list=PLBv09BD7ez_4Bs9j3o8l_ZTjQZoN_3Oqs



Data science is easy
if you know the right
features.

An artificial neural network consists of a number of interconnected processors: **neurons**.

The neurons are connected by **weighted links** passing signals to one another.

Connectionism


Describe complicated phenomena by interconnected networks of simple units.

Self-organization

The structure of a neural network is only partially predetermined.


Neither totally undetermined, nor fully determined!

This is a cute little network designed by me...



This is a poor creature from almost every textbook and blog on neural networks


- organized into layers
- topology is fixed
- all neurons are same





<https://towardsdatascience.com/randomly-wired-neural-networks-92098dbd5175>

<https://stats.stackexchange.com/questions/135035/can-a-neural-network-with-random-connections-still-work-correctly>


Some dogmas here...



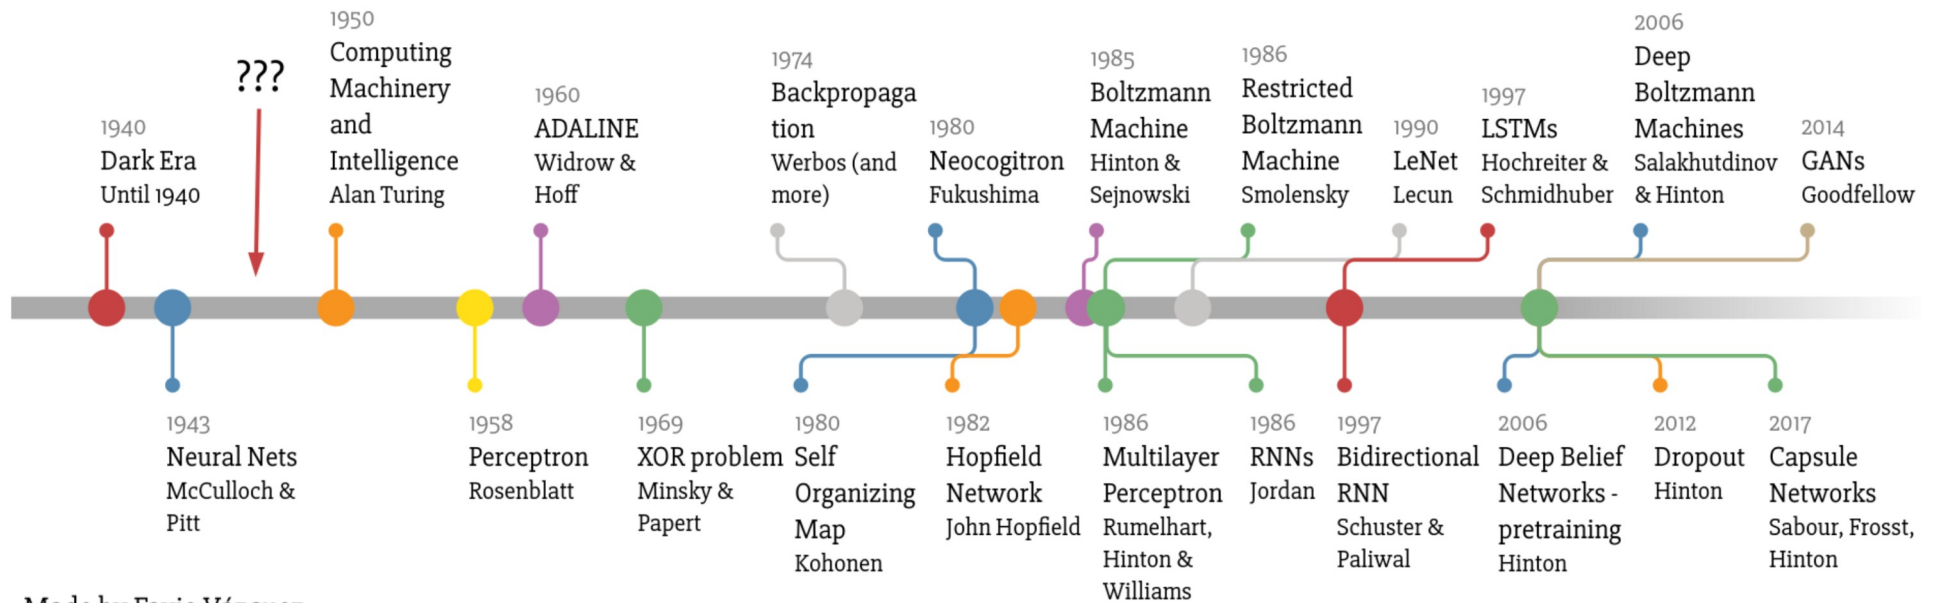
My daddy didn't train neural networks that way,
and his daddy didn't train neural networks that way,
and we don't need no damn physicists coming down here to tell
us how to train neural networks!!!



The classical perceptron [after Rosenblatt 1958]



It is difficult to believe that birds can fly
if you only study ostriches.





Made by Favio Vázquez

<https://towardsdatascience.com/a-weird-introduction-to-deep-learning-7828803693b0>

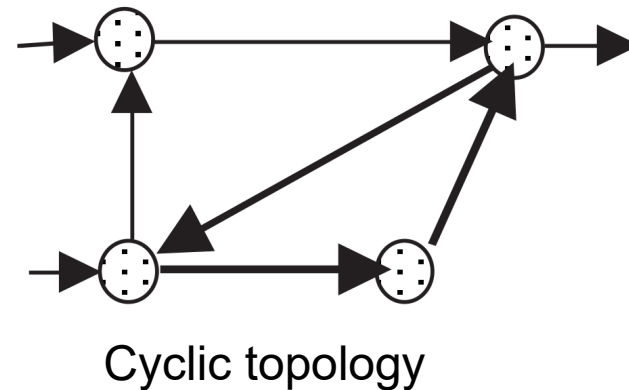
THE MAIN TYPES OF MACHINE LEARNING

https://vas3k.com/blog/machine_learning/





– approximates a nonlinear mapping between its inputs and outputs



aka a **recurrent network** (RNN)

Due to the feedback loop, a recurrent network contains internal memory.

How does a neural network 'learn'

Learning involves modification of the network parameters

- weights

- topology


- or neurons themselves

weights are just regression coefficients

number of neurons
which neurons are connected

functions


by applying a set of training examples.



The training of the network is repeated for many examples in the set until the network reaches a steady state where there are no further significant changes in its parameters.

The previously applied training examples may be reapplied during the training session but in a different order.

Any mathematically defined change in ANN parameters over time is referred to as the **learning algorithm**.




A neuron

- receives input signals
- computes a new activation level
- sends it as an output signal through the output links

The input signal can be raw data or outputs of other neurons.

The output signal can be either a final solution to the problem or an input to other neurons.



$$Y = f(w_1, w_2, \dots, w_n; x_1, x_2, \dots, x_n)$$

f – the **primitive function**
can be any function

Usually, f is a superposition of a **transfer function** and an **activation function**.

$$z = z(w_1, w_2, \dots, w_n; x_1, x_2, \dots, x_n) \quad \text{– transfer function}$$

– aggregates all inputs and weights into a small set of numbers (usually a single number)

$$\text{e.g., } z = \sum_{i=1}^n w_i x_i - \theta \quad \text{which is called a linear combiner.}$$

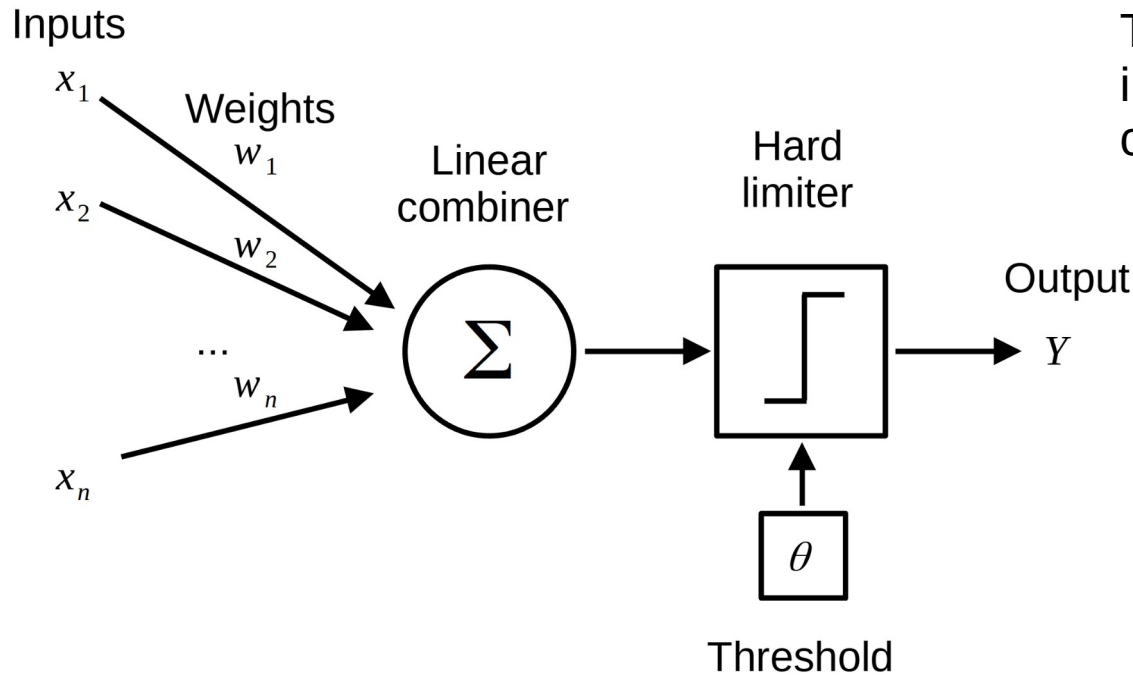
$$Y = \sigma(z) \quad \text{– activation function}$$

– a non-linear transformation of the aggregated input

The transfer and activation functions may introduce some symmetry (desired or unwanted).

The output of the neuron will be the same if $z(w; x) = z(w; \tilde{x})$ or if $\sigma(z) = \sigma(\tilde{z})$



Perceptron



The aim of the perceptron is to classify inputs, x_1, x_2, \dots, x_n , into one of two classes, say A_1 and A_2 .

The n -dimensional input space is divided by a hyperplane into two decision regions.

The hyperplane is defined by the linear function




$$z = \sum_{i=1}^n w_i x_i - \theta$$

The threshold θ shifts the decision boundary.

The **perceptron learning rule**

A perceptron learns by making small adjustments in the weights to reduce the difference between the actual and desired outputs of the perceptron.

The initial weights are randomly assigned, usually in the range $[-0.5, 0.5]$, and then updated to obtain the output consistent with the training examples.



<https://makeyourownneuralnetwork.blogspot.com/2016/01/why-squared-error-cost-function.html>

If $e(p) > 0$, we need to increase perceptron output $Y(p)$
If $e(p) < 0$, we need to decrease $Y(p)$

Each perceptron input contributes $x_i(p)w_i(p)$ to the total input $X(p) \Rightarrow$

If $x_i(p) > 0$, an increase in its weight $w_i(p)$ tends to increase perceptron output $Y(p)$,
whereas if $x_i(p) < 0$, an increase in $w_i(p)$ tends to decrease $Y(p) \Rightarrow$

$$w_i(p+1) = w_i(p) + \alpha x_i(p)e(p)$$

where α is the **learning rate**, a positive constant

$w_i(p+1) = w_i(p) + \alpha x_i(p) e(p)$ – a linear function of error

It can be extended to any non-decreasing function

$$w_i(p+1) = w_i(p) + x_i(p) \sum_{j=0}^{\infty} \alpha_j e(p)^{2j+1}$$

which effectively means that the learning rate is not constant but depends on the error:

$$w_i(p+1) = w_i(p) + \alpha(e(p)) x_i(p) e(p)$$

$$\alpha(e(p)) = \sum_{j=0}^{\infty} \alpha_j e(p)^{2j}$$


where the coefficients α_j are such that α is a non-negative function.

We can also add a more general dependence on x to make the learning rate depend on certain values of the inputs:


$$w_i(p+1) = w_i(p) + \alpha(e(p), x_1(p), \dots, x_n(p)) x_i(p) e(p)$$

where α is a non-negative function


Perceptron performing logical operations



AND ($x_1 \cap x_2$)



OR ($x_1 \cup x_2$)





Exclusive-OR ($x_1 \oplus x_2$)

Separable by a single straight line aka **linearly separable**

Not linearly separable

Either multiple straight lines (multiple neurons) or non-linear transfer function, e.g.



$$z = \sum_{i=1}^n w_i x_i + \sum_{i,j=1}^n w_{ij} x_i x_j - \theta$$

or use different activation function

See e.g. here:

<https://ai.stackexchange.com/questions/9417/why-cant-the-xor-linear-inseparability-problem-be-solved-with-one-perceptron/>

“Usual” activation functions

https://en.wikipedia.org/wiki/Activation_function

Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Logistic (a.k.a. Sigmoid or Soft step)		$f(x) = \sigma(x) = \frac{1}{1 + e^{-x}} \text{ [1]}$
TanH		$f(x) = \tanh(x) = \frac{(e^x - e^{-x})}{(e^x + e^{-x})}$
Rectified linear unit (ReLU) [11]		$f(x) = \begin{cases} 0 & \text{for } x \leq 0 \\ x & \text{for } x > 0 \end{cases} = \max\{0, x\} = x \mathbf{1}_{x>0}$

Layer activations

- relu function
- sigmoid function
- softmax function
- softplus function
- softsign function
- tanh function
- selu function
- elu function
- exponential function

Layer weight initializers

- RandomNormal class
- RandomUniform class
- TruncatedNormal class

usual?

traditional?

venerated?

sacred?

Kolmogorov's theorem
does not say anything
against me!

Sigmoid is not
more non-linear
than I am!

Justice for all activation functions!


Non-traditional activation functions are protesting:

<https://ai.stackexchange.com/questions/24117/smallest-possible-network-to-approximate-the-sin-function>

<https://datascience.stackexchange.com/questions/58838/can-we-learn-fx-1-x-using-a-neural-network-exactly>

<https://stats.stackexchange.com/questions/361066/what-is-the-point-of-having-a-dense-layer-in-a-neural-network-with-no-activation>

Multilayer perceptron



How to update weights when more than one node contributes to an output and its error?

We don't know how much each link contributes to the total error but we can approximate the probability distribution of the error contributions among the links.

The probability that only one link of many was responsible for the error is extremely small.


Links with larger weights contribute more to the error.

Larger output of a neuron \rightarrow larger error

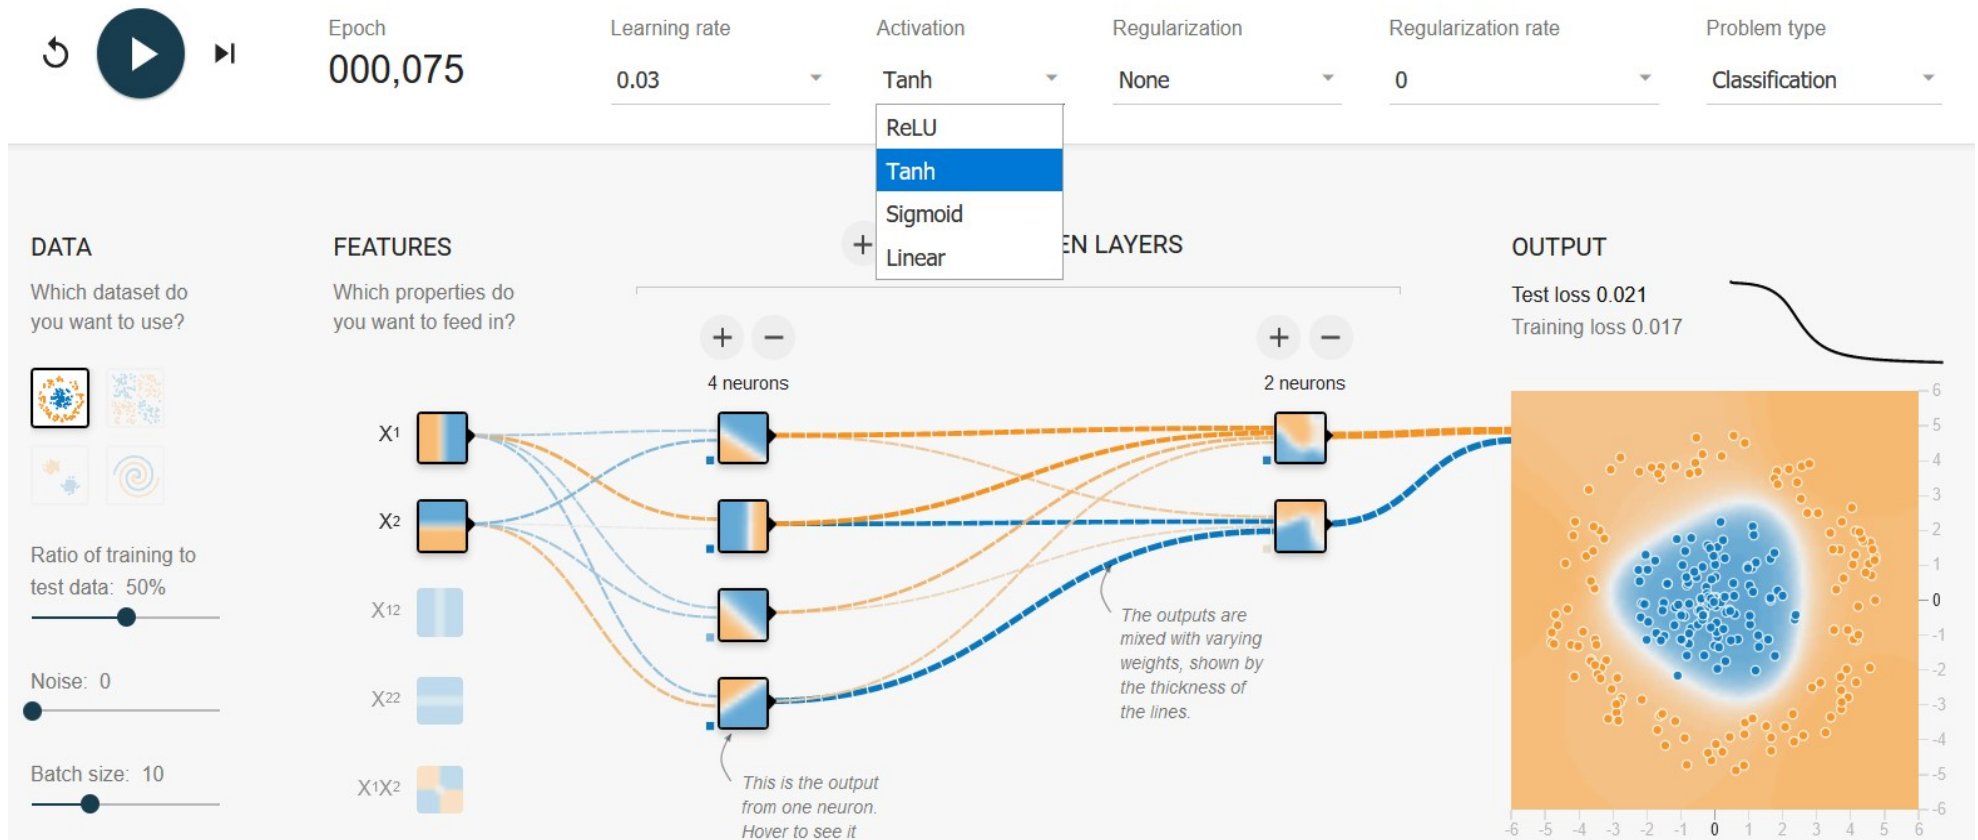
Sharper change through the neuron \rightarrow larger error

The only neuron that never makes a mistake is the neuron who never does anything.

<http://playground.tensorflow.org>



Tinker With a **Neural Network** Right Here in Your Browser.
Don't Worry, You Can't Break It. We Promise.



The back-propagation algorithm: derivation

$$\frac{\partial E}{\partial w_{jk}} = \frac{\partial E}{\partial o_k} \frac{\partial o_k}{\partial w_{jk}}$$

$$\frac{\partial E}{\partial w_{jk}} = \frac{\partial E}{\partial o_k} \frac{\partial \sigma_k(z_k)}{\partial w_{jk}} = \frac{\partial E}{\partial o_k} \sigma'_k \frac{\partial z_k}{\partial w_{jk}}$$

$$\frac{\partial E}{\partial w_{jk}} = \delta_k \frac{\partial z_k}{\partial w_{jk}} \quad \text{where} \quad \delta_k = \sigma'_k \frac{\partial E}{\partial o_k}$$

$$\frac{\partial E}{\partial w_{ij}} = \sum_{k \in K} \frac{\partial E}{\partial o_k} \frac{\partial o_k}{\partial w_{ij}} = \sum_{k \in K} \frac{\partial E}{\partial o_k} \frac{\partial \sigma_k(z_k)}{\partial w_{ij}} = \sum_{k \in K} \frac{\partial E}{\partial o_k} \sigma'_k \frac{\partial z_k}{\partial w_{ij}} = \sum_{k \in K} \delta_k \frac{\partial z_k}{\partial w_{ij}}$$


$$\frac{\partial z_k}{\partial w_{ij}} = \frac{\partial z_k}{\partial o_j} \frac{\partial o_j}{\partial w_{ij}} = \frac{\partial z_k}{\partial o_j} \frac{\partial \sigma_j(z_j)}{\partial w_{ij}} = \frac{\partial z_k}{\partial o_j} \sigma'_j \frac{\partial z_j}{\partial w_{ij}}$$

$$\frac{\partial E}{\partial w_{ij}} = \sum_{k \in K} \delta_k \frac{\partial z_k}{\partial o_j} \sigma'_j \frac{\partial z_j}{\partial w_{ij}} = \delta_j \frac{\partial z_j}{\partial w_{ij}}$$

$$\text{where} \quad \delta_j = \sigma'_j \sum_{k \in K} \delta_k \frac{\partial z_k}{\partial o_j}$$

$$\frac{\partial E}{\partial w_{li}} = \sum_{k \in K} \delta_k \frac{\partial z_k}{\partial w_{li}} = \sum_{k \in K} \delta_k \sum_{j \in J} \frac{\partial z_k}{\partial o_j} \sigma'_j \frac{\partial z_j}{\partial w_{li}} = \sum_{j \in J} \delta_j \frac{\partial z_j}{\partial w_{li}} = \sum_{j \in J} \delta_j \frac{\partial z_j}{\partial o_i} \sigma'_i \frac{\partial z_i}{\partial w_{li}}$$

$$\frac{\partial E}{\partial w_{li}} = \delta_i \frac{\partial z_i}{\partial w_{li}} \quad \text{where} \quad \delta_i = \sigma'_i \sum_{j \in J} \delta_j \frac{\partial z_j}{\partial o_i}$$




$$w_{ab} \rightarrow w_{ab} + \Delta w_{ab}$$

Gradient descent:

$$\Delta w_{ab} = -\eta \frac{\partial E}{\partial w_{ab}}$$

$$E = \sum_{k \in K} E_k$$

$$o_b = \sigma_b(z_b)$$



The back-propagation algorithm: summary

$$E = \sum_{k \in K} E_k$$

$$o_b = \sigma_b(z_b)$$

$$\delta_k = \sigma'_k \frac{\partial E}{\partial o_k}$$

$$\delta_j = \sigma'_j \sum_{k \in K} \delta_k \frac{\partial z_k}{\partial o_j}$$

$$\delta_i = \sigma'_i \sum_{j \in J} \delta_j \frac{\partial z_j}{\partial o_i}$$

$$\frac{\partial E}{\partial w_{jk}} = \delta_k \frac{\partial z_k}{\partial w_{jk}}$$



$$\frac{\partial E}{\partial w_{ij}} = \delta_j \frac{\partial z_j}{\partial w_{ij}}$$

$$\frac{\partial E}{\partial w_{li}} = \delta_i \frac{\partial z_i}{\partial w_{li}}$$

Gradient descent:

$$\Delta w_{ab} = -\eta \frac{\partial E}{\partial w_{ab}}$$

$$w_{ab} \rightarrow w_{ab} + \Delta w_{ab}$$



If the inputs are aggregated via a linear combiner, then

$$z_b = \sum_a w_{ab} o_a + \theta_b \quad \frac{\partial z_b}{\partial w_{ab}} = o_a \quad \frac{\partial z_b}{\partial o_a} = w_{ab} \quad \frac{\partial z_b}{\partial \theta_b} = 1$$


$$\delta_k = \sigma'_k \frac{\partial E}{\partial o_k} \quad \frac{\partial E}{\partial w_{jk}} = o_j \delta_k \quad \frac{\partial E}{\partial \theta_k} = \delta_k$$

$$\delta_j = \sigma'_j \sum_{k \in K} \delta_k w_{jk} \quad \frac{\partial E}{\partial w_{ij}} = o_i \delta_j \quad \frac{\partial E}{\partial \theta_j} = \delta_j$$

$$\delta_i = \sigma'_i \sum_{j \in J} \delta_j w_{ij} \quad \frac{\partial E}{\partial w_{li}} = o_l \delta_i \quad \frac{\partial E}{\partial \theta_i} = \delta_i$$

$$\Delta w_{ab} = -\eta \frac{\partial E}{\partial w_{ab}} \quad w_{ab} \rightarrow w_{ab} + \Delta w_{ab}$$

$$\Delta \theta_b = -\eta \frac{\partial E}{\partial \theta_b} \quad \theta_b \rightarrow \theta_b + \Delta \theta_b$$



$$o_b = \sigma_b(z_b)$$

$$E = \sum_{k \in K} E_k$$

Correction = error * activation slope * prev. layer output

$$\delta_k = \sigma'_k \frac{\partial E}{\partial o_k}$$

$$\frac{\partial E}{\partial w_{jk}} = o_j \delta_k$$

$$\delta_j = \sigma'_j \sum_{k \in K} \delta_k w_{jk}$$

$$\frac{\partial E}{\partial w_{ij}} = o_i \delta_j$$


$$\delta_i = \sigma'_i \sum_{j \in J} \delta_j w_{ij}$$

$$\frac{\partial E}{\partial w_{li}} = o_l \delta_i$$

$$\Delta w_{ab} = -\eta \frac{\partial E}{\partial w_{ab}}$$

$$w_{ab} \rightarrow w_{ab} + \Delta w_{ab}$$

Correction = error * activation slope * prev. layer output



$$o_b = \sigma_b(z_b)$$

$$z_b = \sum_a w_{ab} o_a$$

$$E = \sum_{k \in K} E_k$$

Neural networks are trained in a series of **epochs**.

An epoch is one forward pass and one back-propagation pass over all training samples.

Full batch learning

The average of the gradients of all the training examples is used in order to update the weights.

– we move (almost) directly towards an optimal solution

Online learning

A single example is used to update the weights.

Approximations
to the gradient


Mini-batch learning

The training set is split into mini-batches. An update is made using the mean-gradient of the mini-batch.

<https://ai.stackexchange.com/questions/19894/are-there-any-commonly-used-discontinuous-activation-functions>

<https://ai.stackexchange.com/questions/17609/in-deep-learning-is-it-possible-to-use-discontinuous-activation-functions>

Error minimization is a global optimization problem



Problems:

Local minima (not a big problem for large-size networks)

Anna Choromanska et al. The Loss Surfaces of Multilayer Networks <https://arxiv.org/abs/1412.0233v3>

LeCun, Y., Bengio, Y. & Hinton, G. Deep learning. Nature 521, 436–444 (2015). <https://doi.org/10.1038/nature14539>

Saddle points (a bigger problem)

Yann Dauphin et al., Identifying and attacking the saddle point problem in high-dimensional non-convex optimization <https://arxiv.org/abs/1406.2572>


For fairly small networks

Mavrovouniotis, M., Yang, S. Training neural networks with ant colony optimization algorithms for pattern classification. Soft Comput 19, 1511–1522 (2015). <https://doi.org/10.1007/s00500-014-1334-5>

They show that ant colony optimization + backprop beats unmodified backprop on several benchmark data sets (albeit not by much).

Liao, SH., Hsieh, JG., Chang, JY. et al. Training neural networks via simplified hybrid algorithm mixing Nelder–Mead and particle swarm optimization methods. Soft Comput 19, 679–689 (2015). <https://doi.org/10.1007/s00500-014-1292-y>

Search for the best hyperparameters is another global optimization problem



How many layers?

How many neurons in each layer?


Which activation function?

Learning rate?

Training algorithm?

Batch size?

Any regularization?



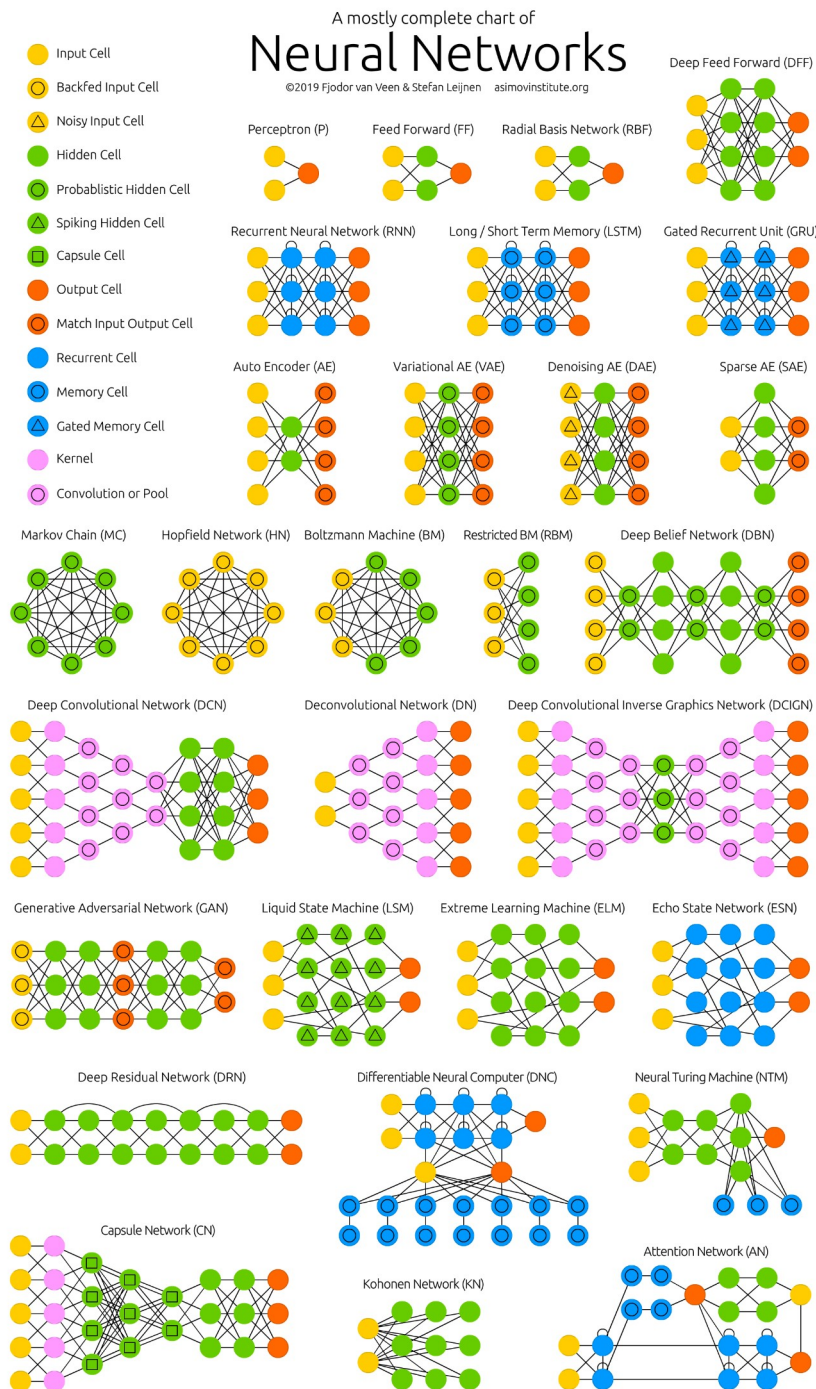
<https://www.coursera.org/learn/neural-networks-deep-learning>

Ozaki, Y., Yano, M. & Onishi, M. Effective hyperparameter optimization using Nelder-Mead method in deep learning. IPSJ T Comput Vis Appl 9, 20 (2017). <https://doi.org/10.1186/s41074-017-0030-7>

Rios, L.M., Sahinidis, N.V. Derivative-free optimization: a review of algorithms and comparison of software implementations. J Glob Optim 56, 1247–1293 (2013). <https://doi.org/10.1007/s10898-012-9951-y>

The concepts behind efficient hyperparameter tuning using Bayesian optimization

<https://towardsdatascience.com/a-conceptual-explanation-of-bayesian-model-based-hyperparameter-optimization-for-machine-learning-b8172278050f>



To solve a problem with machine learning, we need either of the following:

- a good set of features
- a lot of data (millions, billions)
- a good network architecture and learning algorithm

More complex problems require larger networks.

Larger networks contain more parameters.

More parameters require more data.

If we try to fit a large network with too little data, the model will overfit and make worse predictions than a simpler network.

Even a simple regression model can easily beat a large neural network if there is only scarce data.

Links

A Neural Network in Python <http://iamtrask.github.io/2015/07/12/basic-python-network/>
<https://iamtrask.github.io/2015/07/27/python-network-part2/>

A Fortran version: <https://github.com/burubaxair/machine-learning-in-fortran/blob/main/nn.f90>

Grokking-Deep-Learning <https://github.com/iamtrask/Grokking-Deep-Learning>

Backpropagation Video Tutorials

<https://makeyourownneuralnetwork.blogspot.com/2015/04/backpropagation-video-tutorials.html>

Draw network architecture diagrams

<https://datascience.stackexchange.com/questions/14899/how-to-draw-deep-learning-network-architecture-diagrams>

Some “usual” activation functions

<https://stats.stackexchange.com/questions/115258/comprehensive-list-of-activation-functions-in-neural-networks-with-pros-cons>

How to choose the number of hidden layers and neurons

<https://stats.stackexchange.com/questions/181/how-to-choose-the-number-of-hidden-layers-and-nodes-in-a-feedforward-neural-netw>

<https://ai.stackexchange.com/questions/20680/should-neural-nets-be-deeper-the-more-complex-the-learning-problem-is>

<https://datascience.stackexchange.com/questions/26597/how-to-set-the-number-of-neurons-and-layers-in-neural-networks>

How to implement a neural network <https://peterroelants.github.io/posts/neural-network-implementation-part01/>

37 Reasons why your Neural Network is not working

<https://blog.slavv.com/37-reasons-why-your-neural-network-is-not-working-4020854bd607>

Gradient Descent on Riemannian Manifolds

<https://wiseodd.github.io/techblog/2019/02/22/optimization-riemannian-manifolds/>